

REINFORCEMENT LEARNING

Hyundai Motors Boot-campus

UKJO HWANG, DOHYEOK KWON

INFORMATION AND INTELLIGENCE SYSTEMS LAB.

ELECTRONIC ENGINEERING, HANYANG UNIVERSITY


July 7, 2025

GOAL

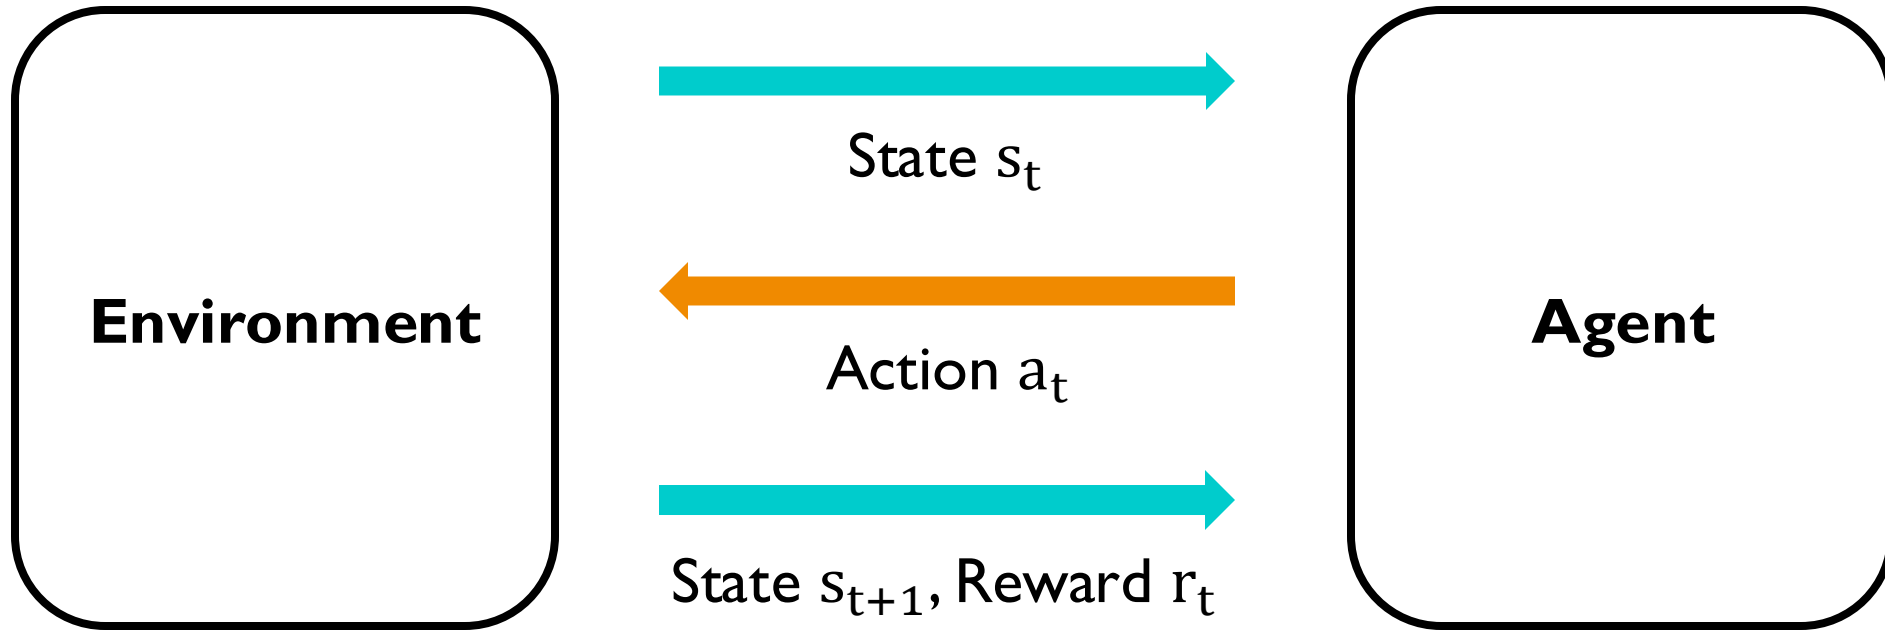
- **How to implement deep RL algorithms**
 - How to utilize environments using the OpenAI Gym API
 - How to train neural networks for reinforcement learning tasks
 - How to train agents using DQN and DDPG algorithms



ENVIRONMENT AGENT INTERACTION



ENVIRONMENT AGENT INTERACTION




- The environment provides the agent with information about the current **state** and **reward**.
- The agent selects an **action** based on its policy.
- The environment transitions to a new state according to its **state transition probabilities**.

SUPERVISED LEARNING VS REINFORCEMENT LEARNING


- **Supervised Learning**

- The data is provided in advance.
- The distribution and quality of data are given beforehand.



- **Reinforcement Learning**

- The data is **generated through interactions** between the agent and the environment during the learning process
- The distribution and quality of data continuously change according to the action (policy).



HOW TO IMPLEMENT RL ALGORITHMS

Environment

- What is the shape and structure of the data?
- What functions and APIs does the environment provide?



Agent

- How to design and implement the model (e.g., Neural Network)
- How to train the model using the provided data



Combination

- How the agent interacts with the environment
- How to implement and apply specific RL algorithms

OPENAI GYMNASIUM

- **OpenAI Gymnasium**

- Provides a wide range of environments for RL.
- It's important to understand the basic interface of Gymnasium.
- To implement a new environment, it's best to follow the Gymnasium API standards.
- <https://gymnasium.farama.org/>



SETUP

- **I. Create Anaconda virtual environment**

- `conda create -n hyundai_rl python=3.10`
- `conda activate hyundai_rl`

- **II. Install packages**

- `pip install gymnasium matplotlib ipykernel torch`



I. Q-LEARNING

BACKGROUND

- **Q-Value Function(or State-Action Value Function) $Q(S, A)$**

- 특정상태 $S(State)$ 에서 특정 행동 $A(Action)$ 를 선택했을 때 기대되는 누적 보상.

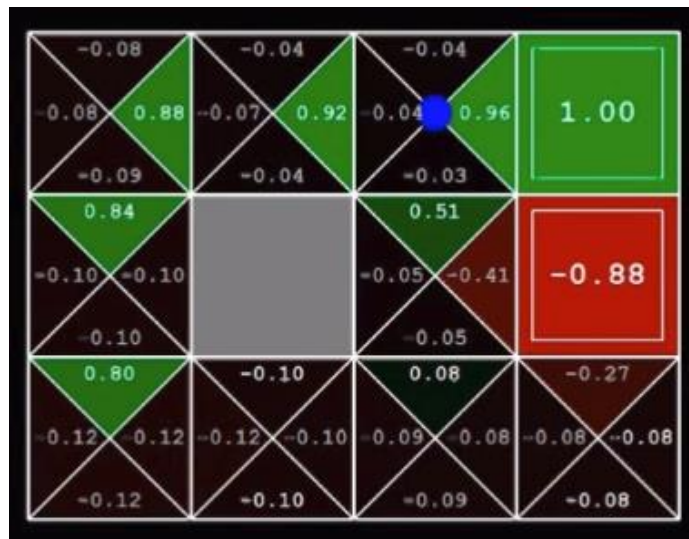

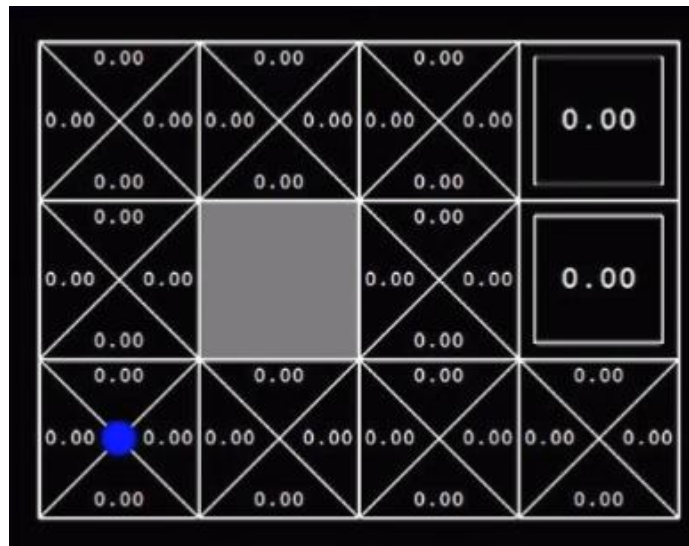
$$Q(S, A) = \mathbb{E}[R_t | s_t = S, a_t = A] = \mathbb{E} \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k} \mid s_t = S, a_t = A \right]$$

- 학습을 통해서 모든 state와 action에 대한 Q 값을 찾으면, 특정 state에서 어떤 action을 선택하는 것이 가장 큰 보상을 받을 수 있을지 아는 것과 같음.

BACKGROUND

- **Q-Table**

- Q-Value를 저장하는 간단한 방법.
- 모든 state, action pair에 대해서 Q-value를 저장해야 함.
- Dimension: $S \times A$



	Left	Right	Up	Down
(1,1)	0	0	0	0
(1,2)	0	0	0	0
(1,3)	0	0	0	0
(2,1)	0	0	0	0
(2,2)	0	0	0	0
(2,3)	0	0	0	0
(3,1)	0	0	0	0
(3,2)	0	0	0	0
(3,3)	0	0	0	0

Q-LEARNING

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

 until S is terminal

Q-LEARNING

- 1. Q-Table 초기화

$$Q(S, A) = 0 \quad \forall S, A$$

- 2. ϵ -greedy policy

- 탐색(Exploration)과 활용(Exploitation)을 적절히 조합하여 다음 action을 선택해야 함.
- ϵ 의 확률로 random action을 선택하고, $1 - \epsilon$ 의 확률로 현재 state에서 가장 높은 Q-value를 갖는 action 선택


- 3. Q-value update

- Bellman equation을 따라 Q-value를 update.

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_{a'} Q(S', a') - Q(S, A)]$$

- α : Learning rate / γ : Discount factor \Rightarrow 모두 $[0, 1]$ 사이의 실수 값

ENV: CLIFF WALKING




- **Action Space:** Discrete(4) | **Observation Space:** Discrete(48)
 - **[Action]** 0: Move up | 1: Move right | 2: Move down | 3: Move left
- **Starting State:** [36] (3, 0) | **Episode End:** [47] (3, 11)
- **Reward**
 - Each time step incurs -1 reward, unless the player stepped into the cliff, which incurs -100 reward.


ENV: CLIFF WALKING

- **[Training]**


- Episode 0~10 (Average steps: 240.0)



- Episode 100~110 (Average steps: 38.5)



- Episode 200~210 (Average steps: 17.1)




EXPERIMENTAL EXTENSIONS

- **Exploration probability** 가 더 높은 값이라면?
 - Increase ‘exploration rate’: 0.1 → **0.4**
 - Decrease ‘number of episodes’: 300 → **150**
- 유사한 환경인 ‘FrozenLake’에서 학습에 **더 많은 step**이 필요한 이유?
 - `env = gym.make('FrozenLake-v1', render_mode='rgb_array', is_slippery=False)`
 - Increase episodes from 150 to 2000
- **Stochastic transition**의 영향 확인
 - Set ‘is_slippery = True’

EXPERIMENTAL EXTENSIONS

- **Exploration probability** 가 더 높은 값이라면?
- 유사한 환경인 'FrozenLake'에서 학습에 **더 많은 step**이 필요한 이유?
- **Stochastic transition**의 영향 확인



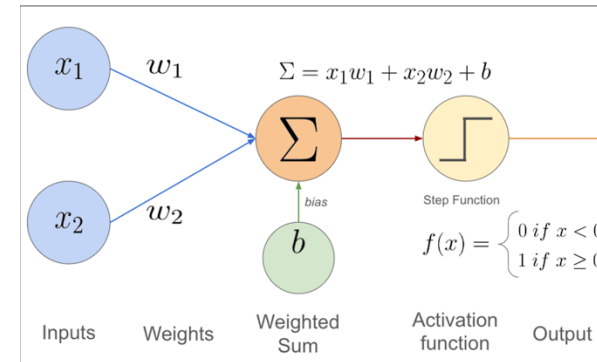

Slippery



II. DQN


DEEP NEURAL NETWORKS

- If the state and action spaces are large, it becomes impractical to implement the Q-function using a table.
- Therefore, deep neural networks are widely used to approximate the Q-function.



Q-FUNCTION


- The output of the DNN is the Q-value corresponding to a given state-action pair.
- Therefore, it is natural to design the DNN to take the state and action as input and output the corresponding Q-value.



Q-FUNCTION


- However, in practice, the DNN is often designed to take only the state as input and output the Q-values for every possible action.
- The feasibility of this implementation relies on the discreteness of the action space.

$$\nabla_{\theta_i} L_i(\theta_i) = \mathbb{E}_{s, a \sim \rho(\cdot); s' \sim \mathcal{E}} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) - Q(s, a; \theta_i) \right) \nabla_{\theta_i} Q(s, a; \theta_i) \right]$$



IMPLEMENTATION DETAILS

- Since DQN is an off-policy algorithm, it utilizes a replay buffer to store samples generated while interacting with the environment.
- This design choice is crucial because it helps to reduce strong correlations between samples.



IMPLEMENTATION DETAILS

- DQN employs a target network that gradually updates to follow the parameters of the original network.
- This mechanism contributes to stabilizing the training.

$$\nabla_{\theta_i} L_i(\theta_i) = \mathbb{E}_{s, a \sim \rho(\cdot); s' \sim \mathcal{E}} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) - Q(s, a; \theta_i) \right) \nabla_{\theta_i} Q(s, a; \theta_i) \right].$$

DQN ALGORITHM

Algorithm 1 Deep Q-learning with Experience Replay

Initialize replay memory \mathcal{D} to capacity N

Initialize action-value function Q with random weights

for episode = 1, M **do**

 Initialise sequence $s_1 = \{x_1\}$ and preprocessed sequenced $\phi_1 = \phi(s_1)$

for $t = 1, T$ **do**

 With probability ϵ select a random action a_t

 otherwise select $a_t = \max_a Q^*(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in \mathcal{D}

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from \mathcal{D}

 Set $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ according to equation 3

end for

end for



III. ACTOR-CRITIC ALGORITHM

BACKGROUND

- **Policy-based reinforcement learning**

- 기존 Value-based algorithm이 어떤 state에서 각 action의 가치를 계산하여 행동했다면, Policy-based algorithm은 각 state에서의 행동 정책(확률)을 학습하는 방식.
 - › DQN은 Network의 output으로 특정 (s, a)의 Q-Value를 출력
 - › Policy Network는 각 현재 state에서 각 action을 수행할 확률 $\pi(a|s) = P(A_t = a | S_t = s)$ 를 출력.

- **Policy gradient**

- Objective function: $J(\theta) = V_{\pi_\theta}(s_0)$ (To maximize it, $\theta_{t+1} = \theta_t + \alpha \nabla_\theta J(\theta)$)
 - › $V_{\pi_\theta}(s) = \mathbb{E}_{\pi_\theta}[R_{t+1} + \gamma V_{\pi_\theta}(S_{t+1}) | S_t = s]$ (Expected return)
- **Policy gradient theorem:** $\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta}[r(\tau) \sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t | s_t)]$


BACKGROUND

- **Policy gradient**

- **Policy gradient theorem:** $\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}}[r(\tau) \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)]$

$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}}[\sum_{t=0}^{T-1} G_t \nabla_{\theta} \log \pi_{\theta}(a_t s_t)]$	REINFORCE (Monte Carlo PG)
$= \mathbb{E}_{\pi_{\theta}}[\sum_{t=0}^{T-1} (G_t - V_{\phi}(s_t)) \nabla_{\theta} \log \pi_{\theta}(a_t s_t)]$	REINFORCE with baseline
$= \mathbb{E}_{\pi_{\theta}}[Q_{\phi}(s, a) \nabla_{\theta} \log \pi_{\theta}(a s)]$	Q-value Actor-Critic
$= \mathbb{E}_{\pi_{\theta}}[A_{\phi_1, \phi_2}(s, a) \nabla_{\theta} \log \pi_{\theta}(a s)]$	Advantage Actor-Critic
$= \mathbb{E}_{\pi_{\theta}}[(r + \gamma V_{\phi}(s') - V_{\phi}(s)) \nabla_{\theta} \log \pi_{\theta}(a s)]$	TD Actor-Critic
Critic (value function) Actor (policy)	

ACTOR-CRITIC



- **Actor (Policy Network)**

- 각 state s 에서 다음 액션에 대한 확률(정책- $\pi(s|a)$)을 근사하는 Network
- Loss: $J(\theta)$

- **Critic (Value Network)**

- 각 state s 에서의 가치(Value function)를 근사하는 Network.
- Loss: MSE between $V(s)$ and TD target $(r + \gamma V(s'))$

```
01 : Input : Initial Actor policy parameters  $\pi(a|s, \theta)$ 
          : Initial Critic V-function parameters  $V(s, \psi)$ 
02 : Parameters : actor learning rate  $\alpha^\theta > 0$ , critic learning rate  $\alpha^\psi > 0$ 
03 : trial_step =  $T$ 
04 : For episode = 1,  $M$  do
05 :     done = False
06 :     Reset environment state.
07 :     For  $t = 1, T$  do (or While not done)
08 :         Observe state  $s$  and select action  $a = \mu_\theta(s)$ ,           # discrete model

09 :         Execute  $a$  in the environment
10 :         Observe next state  $s'$ , reward  $r$ , and done signal  $d$  to indicate whether  $s'$  is terminal
11 :         Reset gradient  $d\theta$  and  $d\psi$  to 0
12 :         Calculate the TD

                                 $TD \leftarrow r + \gamma(1 - d)V(s', \psi) - V(s, \psi)$     (if  $s'$  is terminal, then  $V(s', \psi) \doteq 0$ )
                                where,  $Q_{expected} \leftarrow r + \gamma(1 - d)V(s', \psi)$ 

13 :         Accumulate the policy gradient using the critic:
                                 $d\theta \leftarrow d\theta + \nabla_\theta \log \pi_\theta(s_t, a_t)(TD)$ 


14 :         Accumulate the critic gradient:
                                 $d\psi \leftarrow d\psi + \nabla_\psi (TD)^2$ 

15 :         Update the actor and the critic with the accumulated gradients using gradient descent or similar :

                                 $\theta \leftarrow \theta + \alpha^\theta d\theta$      $\psi \leftarrow \psi + \alpha^\psi d\psi$ 

16 :          $s \leftarrow s'$ 
17 :     End For
18 : End For
19 :
```

ENV: CART POLE




Num	Observation	Min	Max
0	Cart Position	-4.8	4.8
1	Cart Velocity	-Inf	Inf
2	Pole Angle	~ -0.418 rad (-24°)	~ 0.418 rad (24°)
3	Pole Angular Velocity	-Inf	Inf

- **Action Space:** Discrete(2) | **Observation Space:** Continuous(4,)
 - **[Action]** 0: Push cart to the left | 1: Push cart to the right
- **Episode End:** Angle $\pm 12^\circ$ / End of display / 500 steps
- **Reward**
 - A +1 reward is given for every step, including the terminal step.


ENV: CART POLE

- **[Training]**


- Episode 100~200 (Average steps: 60.6)



- Episode 200~210 (Average steps: 144.6)



- Episode 500~510 (Average steps: 282.0)





IV. DDPG

DQN vs DDPG

- DQN

- Designed for environments with discrete action spaces
- A Q-network alone is sufficient for learning the policy

$$\pi(s) = \max_{a \in \mathcal{A}} Q^\pi(s, a)$$

$$Q^*(s, a) = \mathbb{E}_{s' \sim \mathcal{E}} \left[r + \gamma \max_{a'} Q^*(s', a') \mid s, a \right]$$

- DDPG

- Suitable for environments with continuous action spaces
- Requires both a policy network (Actor) and a Q-network (Critic)

$$Q^\mu(s_t, a_t) = \mathbb{E}_{r_t, s_{t+1} \sim E} [r(s_t, a_t) + \gamma Q^\mu(s_{t+1}, \mu(s_{t+1}))]$$

DETERMINISTIC POLICY

- Deterministic policies simplify the computation of gradients during the optimization of the policy network.

$$\begin{aligned}\nabla_{\theta^\mu} J &\approx \mathbb{E}_{s_t \sim \rho^\beta} \left[\nabla_{\theta^\mu} Q(s, a | \theta^Q) \Big|_{s=s_t, a=\mu(s_t | \theta^\mu)} \right] \\ &= \mathbb{E}_{s_t \sim \rho^\beta} \left[\nabla_a Q(s, a | \theta^Q) \Big|_{s=s_t, a=\mu(s_t)} \nabla_{\theta_\mu} \mu(s | \theta^\mu) \Big|_{s=s_t} \right]\end{aligned}$$

CRITIC

- The critic network receives the state and action as input and produces the corresponding Q-value as output.




$$L(\theta^Q) = \mathbb{E}_{s_t \sim \rho^\beta, a_t \sim \beta, r_t \sim E} \left[(Q(s_t, a_t | \theta^Q) - y_t)^2 \right]$$

$$y_t = r(s_t, a_t) + \gamma Q(s_{t+1}, \mu(s_{t+1}) | \theta^Q).$$

ACTOR

- The actor receives the state as input and produces the appropriate action for the agent.

$$\begin{aligned}\nabla_{\theta^\mu} J &\approx \mathbb{E}_{s_t \sim \rho^\beta} [\nabla_{\theta^\mu} Q(s, a | \theta^Q) |_{s=s_t, a=\mu(s_t | \theta^\mu)}] \\ &= \mathbb{E}_{s_t \sim \rho^\beta} [\nabla_a Q(s, a | \theta^Q) |_{s=s_t, a=\mu(s_t)} \nabla_{\theta^\mu} \mu(s | \theta^\mu) |_{s=s_t}]\end{aligned}$$

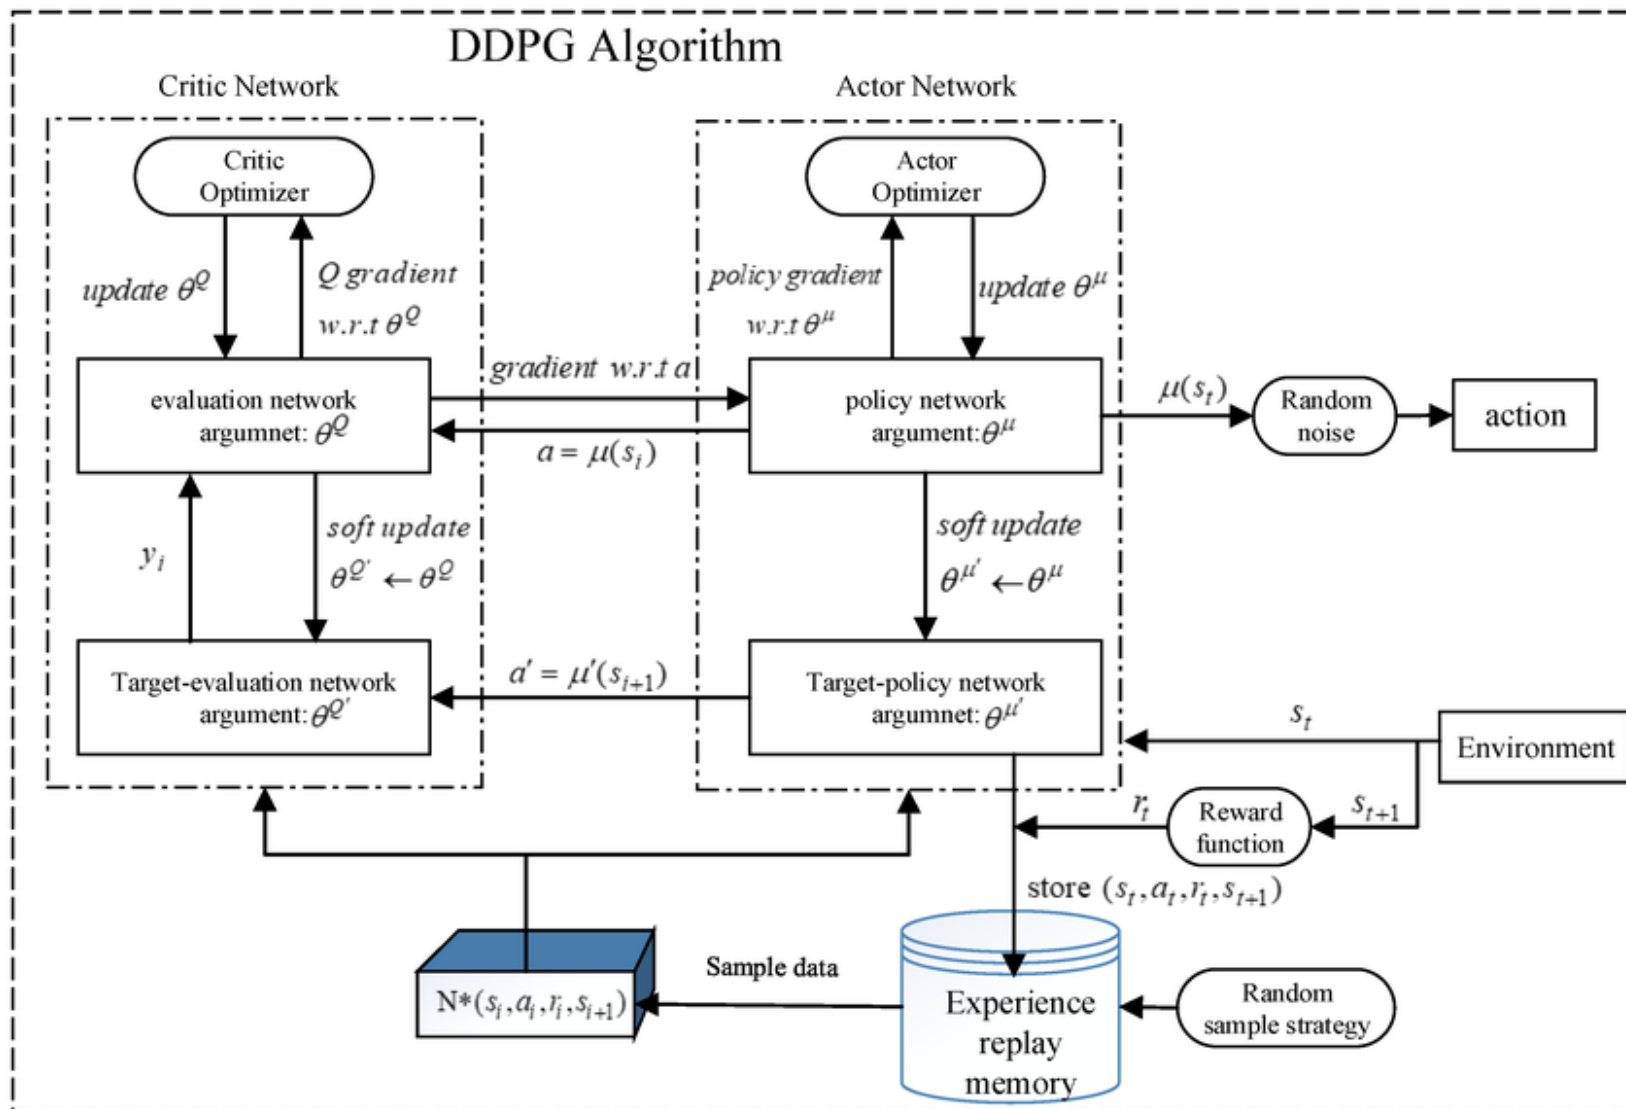


IMPLEMENTATION DETAILS

- Due to the deterministic nature of the policy, the agent may fail to explore enough, which can hinder effective training.
- Therefore, a noise component is incorporated into the agent's policy during training, allowing the agent to explore the environment more effectively.

$$a_t = \mu(s_t | \theta^\mu) + \mathcal{N}_t$$

DDPG STRUCTURE



DDPG ALGORITHM

Algorithm 1 DDPG algorithm

Randomly initialize critic network $Q(s, a|\theta^Q)$ and actor $\mu(s|\theta^\mu)$ with weights θ^Q and θ^μ .

Initialize target network Q' and μ' with weights $\theta^{Q'} \leftarrow \theta^Q, \theta^{\mu'} \leftarrow \theta^\mu$

Initialize replay buffer R

for episode = 1, M **do**

 Initialize a random process \mathcal{N} for action exploration

 Receive initial observation state s_1

for t = 1, T **do**

 Select action $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$ according to the current policy and exploration noise

 Execute action a_t and observe reward r_t and observe new state s_{t+1}

 Store transition (s_t, a_t, r_t, s_{t+1}) in R

 Sample a random minibatch of N transitions (s_i, a_i, r_i, s_{i+1}) from R

 Set $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'})|\theta^{Q'})$

 Update critic by minimizing the loss: $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$

 Update the actor policy using the sampled policy gradient:

$$\nabla_{\theta^\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

 Update the target networks:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$

end for

end for
